

Overview

- Introduction – Linux philosophy
- Command line basics – getting help
- The shell revisited: some features

Navigating the file system

File manipulation

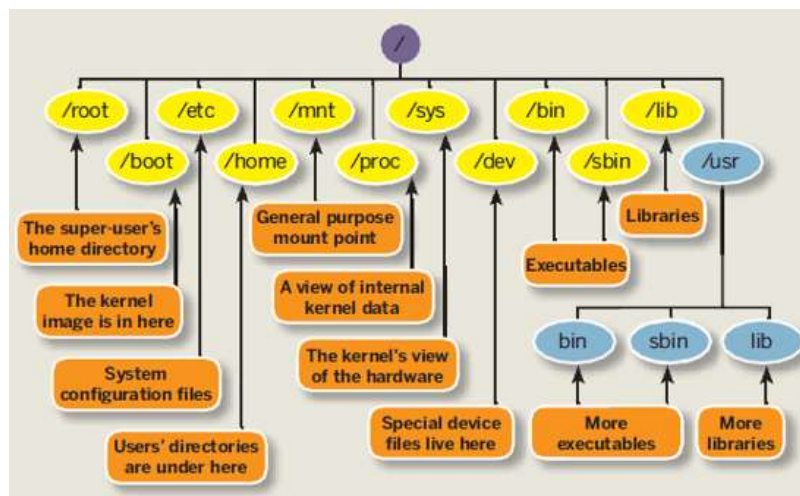
- Text editing
- Various commands
- Archiving
- Groups, users, security
- Process control

Linux filesystem
Navigating the filesystem

Linux File System

- *hierarchical directory structure*: files are organized in a tree-like pattern of directories (folders), which may contain files and other directories, etc.
- Everything is a file:
 - Regular files
 - Directories: files listing a set of files
 - Symbolic links: files referring to the name of another file
- *root /*: the first directory in the file system.
- Note: comparison with Windows,
 - Windows has a separate file system tree for each storage device (e.g. C-drive, D-drive, I-drive, ...)
 - Linux has a single file system tree, regardless of how many drives or storage devices are attached to the computer.
Storage devices are attached (or *mounted*) at various points on the tree.

Linux File System



Source: <http://linuxsuperuser07.blogspot.be/2011/09/rhel-6-file-system.html>

Linux File System

- Not imposed by the system. Can vary from one system to the other, even between two GNU/Linux installations!
- The Unix filesystem structure is defined by the Filesystem Hierarchy Standard (FHS):
<http://www.pathname.com/fhs/pub/fhs-2.3.html>

/	Root Everything is located under /
/bin	Essential User Binaries essential user binaries (programs) that must be present
/boot	Static boot files files needed to boot the system
/dev	Device files special files that represent devices
/etc	Configuration Files contains system-wide configuration files

<https://www.howtogeek.com/117435/htg-explains-the-linux-directory-structure-explained/>

Linux File System

/home	Home folders a home folder for each user.
/lib	Essential Shared Libraries libraries needed by the essential binaries in the /bin and /sbin folder.
/media	Removable Media subdirectories where removable media devices inserted into the computer are mounted.
/mnt	Temporary Mount Points
/opt	Optional Packages optional software packages. It's commonly used by proprietary software.
/proc	Kernel & Process Files
/root	Root Home Directory home directory of the root user

<https://www.howtogeek.com/117435/htg-explains-the-linux-directory-structure-explained/>

Linux File System

/run	Application State Files
/sbin	System Administration Binaries similar to the /bin directory. It contains essential binaries that are generally intended to be run by the root user for system administration.
/srv	Service Data data for services provided by the system
/tmp	Temporary Files Applications store temporary files in the /tmp directory. These files are generally deleted whenever your system is restarted and may be deleted at any time
/usr	User Binaries & Read-Only Data contains applications and files used by users, as opposed to applications and files used by the system. Must be read-only in normal operation.
/var	Variable Data Files the writable counterpart to the /usr directory. Log files and everything else that would normally be written to /usr during normal operation are written to the /var directory.

<https://www.howtogeek.com/117435/htg-explains-the-linux-directory-structure-explained/>

Linux File System

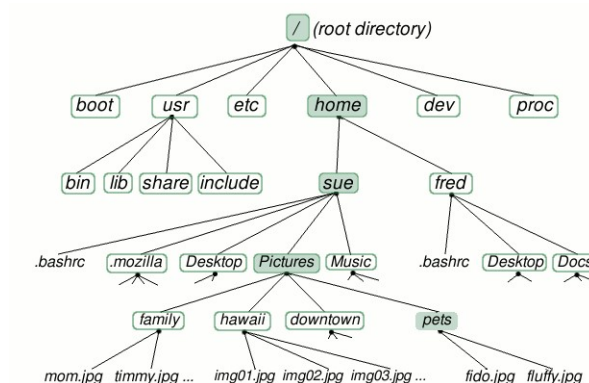
- A **file** is a collection of data, with a location in the file system called a **path**. Paths will typically be a series of words (directory names) separated by forward slashes, /. Files are generally created by users via text editors, compilers, or other means.
- A **directory** is a special type of file. Linux uses a directory to hold information about other files, the equivalent of a folder in Windows. You can think of a directory as a container that holds other files or directories.
- A file is typically stored on physical storage media such as a disk (hard drive, flash disk, etc.).
- Every file must have a name because the operating system identifies files by their name.
 - File names may contain any characters, although some special characters (such as spaces, quotes, and parenthesis) can make it difficult to access the file, so you should avoid them in filenames.
 - File names can be as long as 255 characters, so use descriptive names.
 - File names are case sensitive,
 - A hidden file is any file that begins with a "." (not seen with the bare `ls`)

<https://cww.cac.cornell.edu/Linux/files>

Linux File System

- “\” vs. “/”:
 - In Linux, the “/” is the directory separator, and the “\” is an escape character.
 - In Windows, the forward-slash “/” is the command argument delimiter, while the backslash “\” is a directory separator
- Filenames:
 - In Linux, there is no such thing as a file extension. Periods can be placed at any part of the filename, and “extensions” may be interpreted differently by all programs, or not at all.
 - Check with `file` (determines the file type of a file)
 - Windows uses the “.extension” filename convention, (e.g. FILENAME.TXT).

Linux File System-home directory



Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

ls command

Lists the files in the current directory, in alphanumeric order, except files starting with the "." character.

- `$ ls -a` (all)
Lists all the files (including .* files)
- `$ ls -l` (long)
Long listing (type, date, size, owner, permissions)
- `$ ls -t` (time)
Lists the most recent files first
- `$ ls -S` (size)
Lists the biggest files first
- `$ ls -r` (reverse)
Reverses the sort order
- `$ ls -ltr` (options can be combined)
Long listing, most recent files at the end

ls command

- `$ ls *txt`
The shell first replaces *txt by all the file and directory names ending by txt (including .txt), except those starting with ., and then executes the ls command line.
- `$ ls -d .*`
Lists all the files and directories starting with .
-d list directories themselves, not their contents.
`$ ls -d */`
- `$ ls ?.log`
Lists all the files which names start by 1 character and end by .log

<https://www.thegeekstuff.com/2009/07/linux-ls-command-examples/>

ls command

- There are many color codes, but you can often find only the below 7 colors.
 - White (No color code) Regular File or Normal File
 - Blue: Directory
 - Bright Green: Executable File
 - Bright Red: Archive file or Compressed File
 - Magenta: Image File
 - Cyan: Audio File
 - Sky Blue: Symbolic Link File

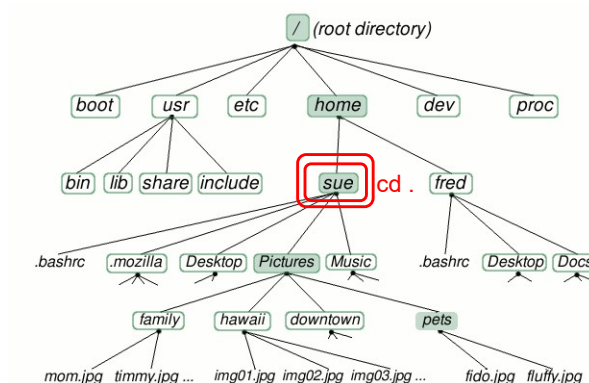
Moving around

- Display the current/working directory
 - `$ pwd`
 - **P**rint **W**orking **D**irectory
 - displays your current location within the file system.
- Change (navigate) directories.
 - `$ cd dir_name`
 - **C**hange **D**irectories
 - changes the position to the specific directory
- You can specify directory names in two ways:
 - Absolute pathname (starts from the root of the tree)
`$ cd /u/home/hpc/test/bin`
 - Relative pathname (relative to your current directory)
`$ cd`
`$ cd .`
`$ cd ..`
`$ cd test/bin`

Special directories

- a few characters representing shortcuts to locations.
- **(single dot) .**
 - The current working directory.
 - Useful to run commands in the current directory
 - `./readme.txt` and `readme.txt` are equivalent.
- **(double dot) ..**
 - The parent (enclosing) directory. Always belongs to the `.` Directory
 - Typical usage:
`cd ..`
- **(tilde) ~**
 - Shells just substitute it by the home directory of the current user.
- **(dash) -**
 - `cd -` – jump back to the previous directory

Linux File System - directories



Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System - directories

```

graph TD
    root["/ (root directory)"] --> boot
    root --> usr
    root --> etc
    root --> home
    root --> dev
    root --> proc
    home --> bin
    home --> lib
    home --> share
    home --> include
    home --> sue
    home --> cd_dot["cd .."]
    home --> fred
    sue --> bashrc1[".bashrc"]
    sue --> mozilla[".mozilla"]
    sue --> Desktop1["Desktop"]
    sue --> Pictures
    sue --> Music
    sue --> bashrc2[".bashrc"]
    sue --> Desktop2["Desktop"]
    sue --> Docs
    Pictures --> family
    Pictures --> hawaii
    Pictures --> downtown
    Pictures --> pets
    family --> mom_jpg["mom.jpg"]
    family --> timmy_jpg["timmy.jpg"]
    family --> dots1["..."]
    hawaii --> img01_jpg["img01.jpg"]
    hawaii --> img02_jpg["img02.jpg"]
    hawaii --> img03_jpg["img03.jpg"]
    hawaii --> dots2["..."]
    pets --> fido_jpg["fido.jpg"]
    pets --> fluffy_jpg["fluffy.jpg"]
  
```

Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>



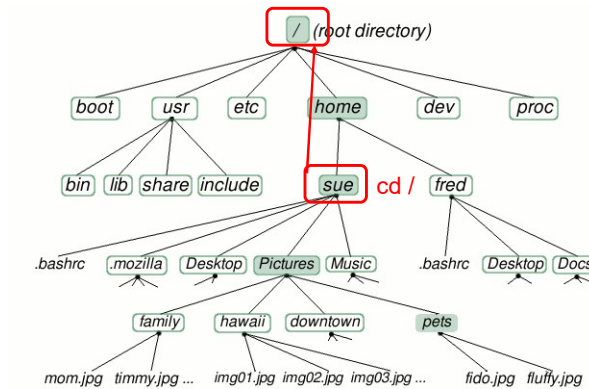
Linux File System - directories

```
graph TD; root["/ (root directory)"] --> boot; root --> usr; root --> etc; root --> home; root --> dev; root --> proc; usr --> bin; usr --> lib; usr --> share; usr --> include; home --> sue; home --> fred; sue --> sue_bashrc[.bashrc]; sue --> sue_mozilla[.mozilla]; sue --> sue_desktop[Desktop]; sue --> sue_pictures[Pictures]; sue --> sue_music[Music]; fred --> fred_bashrc[.bashrc]; fred --> fred_desktop[Desktop]; fred --> fred_docs[Docs]; sue_pictures --> family; sue_pictures --> hawaii; sue_pictures --> downtown; sue_pictures --> pets; mom_jpg[mom.jpg] --> family; timmy_jpg[timmy.jpg] --> family; img01_jpg[img01.jpg] --> hawaii; img02_jpg[img02.jpg] --> hawaii; img03_jpg[img03.jpg] --> hawaii; fido_jpg[fido.jpg] --> pets; fluffy_jpg[fluffy.jpg] --> pets; cd_dot_dot[cd ..] --> pets; style sue stroke:#f00,stroke-width:2px; style pets stroke:#f00,stroke-width:2px; linkStyle 10 stroke:#f00,stroke-width:2px;
```

Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

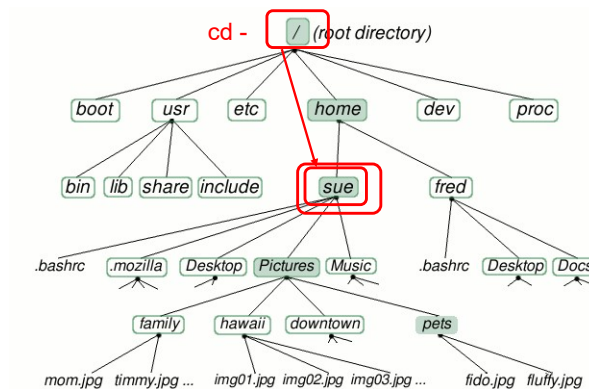


Linux File System - directories



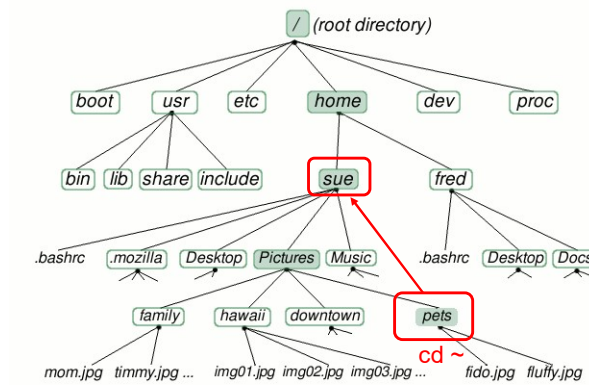
Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System - directories



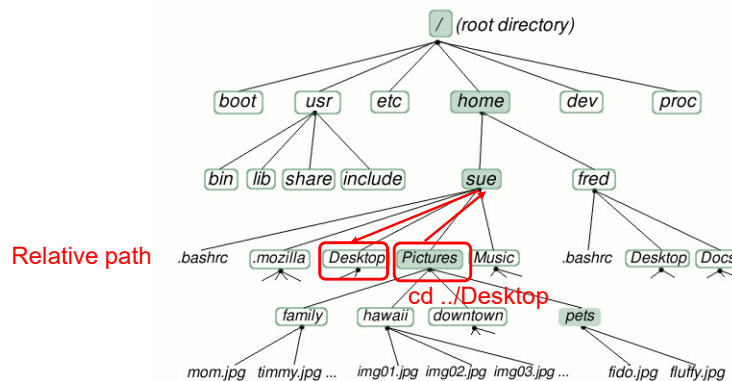
Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System - directories



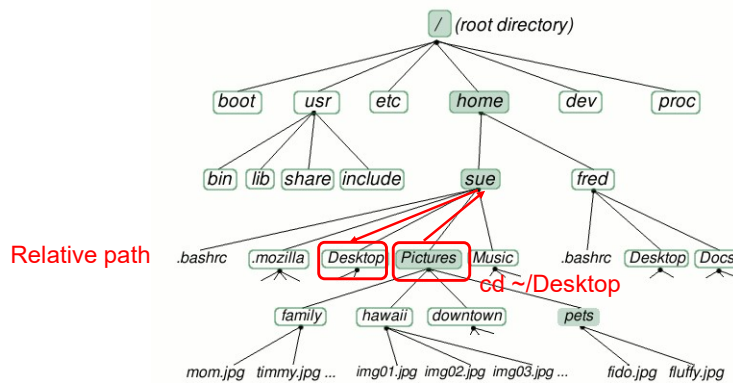
Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System-home directory



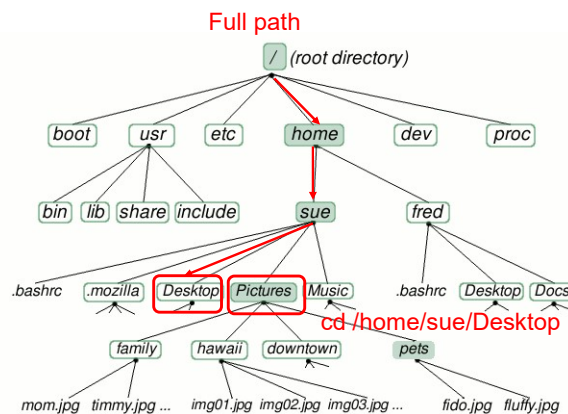
Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System - directories



Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

Linux File System - directories



Source: <http://www.linuxplanet.com/linuxplanet/tutorials/6666/1/screenshot3894/>

File paths

- A path is a sequence of nested directories with a file or directory at the end, separated by the / character
- **Relative path:** `documents/fun/file1`
Relative to the current directory
- **Absolute path:** `/home/user/leuven/file2`
- `/` : root directory.
Start of absolute paths for all files on the system

File manipulation

File Manipulation

- For all file manipulation commands relative or absolute paths can be used (or just files in the current directory when no extra path specified)
- Most of the commands are intuitive – shortcuts of English names

Directories

- Create directories
 - The `mkdir` command is used to create directories
 - `$ mkdir dir1 dir2 dir3`
- Remove directories
 - The `rmdir` command removes directories
 - `$ rmdir dir1`
 - `rmdir` will only remove empty directories.
To remove a non-empty directory, use
`$ rm -r [DIRECTORY]` instead. (BE CAREFUL!)

Copy a file

- The `cp` command copies files and directories
- The default behavior will overwrite any existing file(s). The `-i` option overrides this behavior and prompts the user before overwriting the destination file.
- **syntax:** `cp [OPTIONS] [SOURCE] [DESTINATION]`
 - `$ cp <source_file> <target_file>`
Copies the source file to the target.
 - `$ cp file1 file2 file3 ... dir`
Copies the files to the target directory (last argument).
 - `$ cp -i (interactive)`
Asks for user confirmation if the target file already exists
 - `$ cp -r <source_dir> <target_dir> (recursive)`
Copies the whole directory.
 - `$ cp -v (verbose)`
Displays what has been copied

Move or rename files

- Move or rename files and directories: `mv`
- The default behavior will overwrite any existing file(s).
- **syntax:** `mv [OPTIONS] [SOURCE] [DESTINATION]`
 - `$ mv old_name new_name`
Renames the given file or directory.
 - `$ mv -i (interactive)`
If the new file already exists, asks for user confirm
- The `mv` command can also be used to move or rename directories
 - `$ mv NewFiles/ OldFiles/`
 - `-r` option is not necessary

Remove files

- The `rm` command removes files.
- `$ rm file1 file2 file3 ...`
Removes the given files.
- `$ rm -i` (interactive)
Always ask for user confirmation.
- `$ rm -r dir1 dir2 dir3` (recursive)
Removes the given directories with all their contents.



- Tip:
Whenever you use wildcards with `rm` (besides carefully checking your typing!), test the wildcard first with `ls`. This will let you see the files that will be deleted. Then press the up arrow key to recall the command and replace the `ls` with `rm`.

inode

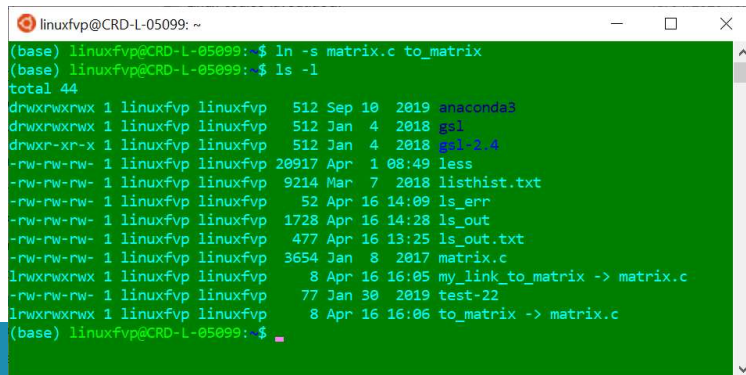
- An inode is a data structure
- Every Linux file or directory has an inode, containing all of the file's metadata
 - a list of all the blocks in which a file is stored,
 - the owner information for that file,
 - permissions and all other attributes that are set for the file
- Check it with
 - `ls -li`

Create links

- **Soft link:** similar to a *shortcut* in Windows. It is an indirect pointer to a file or directory; can point to a file or a directory on a different filesystem or partition.
- Symbolic links are created when using the `-s` option with the `ln` command.

```
ln -s [OPTIONS] FILE LINK
```

- Check with `ls -l`
- The first character “l”, indicates that the file is a symlink.
- The “->” symbol shows the file the symlink points to.



```
linuxfvp@CRD-L-05099: ~  
(base) linuxfvp@CRD-L-05099:~$ ln -s matrix.c to_matrix  
(base) linuxfvp@CRD-L-05099:~$ ls -l  
total 44  
drwxrwxrwx 1 linuxfvp linuxfvp 512 Sep 10 2019 anaconda3  
drwxrwxrwx 1 linuxfvp linuxfvp 512 Jan 4 2018 gsl  
drwxr-xr-x 1 linuxfvp linuxfvp 512 Jan 4 2018 gil-2.4  
-rw-rw-rw- 1 linuxfvp linuxfvp 20917 Apr 1 08:49 less  
-rw-rw-rw- 1 linuxfvp linuxfvp 9214 Mar 7 2018 listhist.txt  
-rw-rw-rw- 1 linuxfvp linuxfvp 52 Apr 16 14:09 ls_err  
-rw-rw-rw- 1 linuxfvp linuxfvp 1728 Apr 16 14:28 ls_out  
-rw-rw-rw- 1 linuxfvp linuxfvp 477 Apr 16 13:25 ls_out.txt  
-rw-rw-rw- 1 linuxfvp linuxfvp 3654 Jan 8 2017 matrix.c  
lrwxrwxrwx 1 linuxfvp linuxfvp 8 Apr 16 16:05 my_link_to_matrix -> matrix.c  
-rw-rw-rw- 1 linuxfvp linuxfvp 77 Jan 30 2019 test-22  
lrwxrwxrwx 1 linuxfvp linuxfvp 8 Apr 16 16:06 to_matrix -> matrix.c  
(base) linuxfvp@CRD-L-05099:~$
```

Create links

- Editing a symbolic link file is the same as editing the source file
- Deleting the symbolic link does not delete the source file.
- Deleting the source file leaves a dangling link
- `$ ln -s file_v5.doc file_final.doc`
creates a symbolic link called `file_final.doc` that points to `file_v5.doc`
- `$ ln -s /home/demo/dir1/dir2/dir3 /home/demo/jump2dir`
creates a symbolic link called `jump2dir` that points to a deep directory - allowing for quicker access

More on files

- Search for files and directories
 - The **find** command performs a raw search on a file system to locate the specified items.
 - ```
$ find location -name some-name
```

```
($ find / -name matrix.c)
```
  - You can also specify more than one location to search,
- Search the locate database for files and directories
  - The **locate** command displays the location of files that match the specified name.
  - Faster than find but lacks the ability to search for advanced characteristics such as file owner, size, and modification time.

## More on files

- Display extended information about a file system, file, or directory
- What does a file contain?
  - Determine a file's type: **file**
  - will print a brief description of the file's contents
  - ```
$ file filename
```
- The **stat** command displays extended information about files. It includes helpful information not available when using the **ls** command

Hands-on 3